

1. Reva Dira Putri (E42241593) Gol B
2. Maria Maksanlina M (E42240119) Gol A
3. Mutiara Ramadhani (E42241957) Gol B
4. Revinka Yos Sheyla (E42241640) Gol B

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import files
uploaded = files.upload()
```

Choose Files | diabetes.csv
diabetes.csv(text/csv) - 23873 bytes, last modified: 11/3/2025 - 100% done
 Saving diabetes.csv to diabetes (2).csv

```
import pandas as pd
data = pd.read_csv('diabetes.csv')
data.head()
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome | grid icon | more icon |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-------|---------|-----------|-----------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | | 0.627 | 50 | 1 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | | 0.351 | 31 | 0 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | | 0.672 | 32 | 1 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | | 0.167 | 21 | 0 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | | 2.288 | 33 | 1 | |

Next steps: [Generate code with data](#) [New interactive sheet](#)

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Pregnancies      768 non-null    int64  
 1   Glucose          768 non-null    int64  
 2   BloodPressure    768 non-null    int64  
 3   SkinThickness    768 non-null    int64  
 4   Insulin          768 non-null    int64  
 5   BMI              768 non-null    float64 
 6   DiabetesPedigreeFunction 768 non-null    float64 
 7   Age              768 non-null    int64  
 8   Outcome          768 non-null    int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
data['Outcome'].value_counts()
```

| count | |
|---------|-----|
| Outcome | |
| 0 | 500 |
| 1 | 268 |


dtype: int64

```
correlation = data.corr()
correlation['Outcome'].sort_values(ascending=False)
```

| | Outcome |
|---------------------------------|----------|
| Outcome | 1.000000 |
| Glucose | 0.466581 |
| BMI | 0.292695 |
| Age | 0.238356 |
| Pregnancies | 0.221898 |
| DiabetesPedigreeFunction | 0.173844 |
| Insulin | 0.130548 |
| SkinThickness | 0.074752 |
| BloodPressure | 0.065068 |

dtype: float64

```
plt.figure(figsize=(10,8))
plt.title('Correlation of Attributes with Outcome variable')
a = sns.heatmap(correlation, square=True, annot=True, fmt='.2f', linecolor='white')
a.set_xticklabels(a.get_xticklabels(), rotation=90)
a.set_yticklabels(a.get_yticklabels(), rotation=30)
plt.show()
```



```
X = data.drop(['Outcome'], axis=1)
y = data['Outcome']
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
X
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | |
|-----|-------------|-----------|---------------|---------------|-----------|-----------|--------------------------|-----------|-----------|
| 0 | 0.639947 | 0.848324 | 0.149641 | 0.907270 | -0.692891 | 0.204013 | | 0.468492 | 1.425995 |
| 1 | -0.844885 | -1.123396 | -0.160546 | 0.530902 | -0.692891 | -0.684422 | | -0.365061 | -0.190672 |
| 2 | 1.233880 | 1.943724 | -0.263941 | -1.288212 | -0.692891 | -1.103255 | | 0.604397 | -0.105584 |
| 3 | -0.844885 | -0.998208 | -0.160546 | 0.154533 | 0.123302 | -0.494043 | | -0.920763 | -1.041549 |
| 4 | -1.141852 | 0.504055 | -1.504687 | 0.907270 | 0.765836 | 1.409746 | | 5.484909 | -0.020496 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 1.827813 | -0.622642 | 0.356432 | 1.722735 | 0.870031 | 0.115169 | | -0.908682 | 2.532136 |
| 764 | -0.547919 | 0.034598 | 0.046245 | 0.405445 | -0.692891 | 0.610154 | | -0.398282 | -0.531023 |
| 765 | 0.342981 | 0.003301 | 0.149641 | 0.154533 | 0.279594 | -0.735190 | | -0.685193 | -0.275760 |
| 766 | -0.844885 | 0.159787 | -0.470732 | -1.288212 | -0.692891 | -0.240205 | | -0.371101 | 1.170732 |
| 767 | -0.844885 | -0.873019 | 0.046245 | 0.656358 | -0.692891 | -0.202129 | | -0.473785 | -0.871374 |

768 rows × 8 columns

Next steps: [Generate code with X](#) [New interactive sheet](#)

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.15, random_state = 0)
X_train.shape, X_test.shape
```

((652, 8), (116, 8))

X_train

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | |
|-----|-------------|-----------|---------------|---------------|-----------|-----------|--------------------------|-----------|-----------|
| 279 | -0.547919 | -0.403562 | -0.367337 | -0.660932 | 1.720954 | -0.849417 | | 1.235602 | -0.956462 |
| 258 | -0.844885 | 2.256695 | -0.987710 | -0.284563 | 2.563195 | -0.773265 | | 0.553055 | -0.786286 |
| 249 | -0.844885 | -0.309671 | 0.873409 | -0.096379 | -0.692891 | -0.240205 | | -0.993245 | -0.871374 |
| 740 | 2.124780 | -0.027996 | 0.563223 | 1.032726 | 0.609544 | 1.308210 | | 0.945671 | 1.255820 |
| 725 | 0.046014 | -0.278373 | 0.459827 | 1.220910 | -0.692891 | 0.940144 | | -0.712374 | 0.404942 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 1.827813 | -0.622642 | 0.356432 | 1.722735 | 0.870031 | 0.115169 | | -0.908682 | 2.532136 |
| 192 | 0.936914 | 1.192592 | -0.160546 | -1.288212 | -0.692891 | -0.202129 | | -0.268417 | 0.234767 |
| 629 | 0.046014 | -0.841722 | -0.212243 | 0.091805 | -0.692891 | -0.925569 | | -0.978145 | -1.041549 |
| 559 | 2.124780 | -1.123396 | 0.253036 | -1.288212 | -0.692891 | -0.240205 | | -0.519087 | 0.149679 |
| 684 | 0.342981 | 0.472758 | 0.666618 | -1.288212 | -0.692891 | -4.060474 | | 0.507754 | 3.042663 |

652 rows × 8 columns

Next steps: [Generate code with X_train](#) [New interactive sheet](#)**X_test**

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | |
|-----|-------------|-----------|---------------|---------------|-----------|-----------|--------------------------|-----------|-----------|
| 661 | -0.844885 | 2.444478 | 0.356432 | 1.409094 | -0.692891 | 1.384362 | | 2.784923 | -0.956462 |
| 122 | -0.547919 | -0.434859 | 0.253036 | 0.593630 | 0.175399 | 0.204013 | | -0.204994 | -0.871374 |
| 113 | 0.046014 | -1.405071 | -0.367337 | -1.288212 | -0.692891 | 0.254780 | | -0.244256 | -0.701198 |
| 14 | 0.342981 | 1.411672 | 0.149641 | -0.096379 | 0.826616 | -0.785957 | | 0.347687 | 1.511083 |
| 529 | -1.141852 | -0.309671 | -0.212243 | -1.288212 | -0.692891 | -0.938260 | | 0.568156 | -0.190672 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 214 | 1.530847 | -0.278373 | 0.666618 | 0.719086 | 0.826616 | 0.280164 | | -0.639892 | 0.234767 |
| 586 | 1.233880 | 0.691838 | -0.160546 | -1.288212 | -0.692891 | 0.369008 | | -1.035527 | 0.660206 |
| 187 | -0.844885 | 0.222381 | 1.493782 | 1.283638 | -0.189283 | 0.000942 | | 2.564454 | -0.020496 |
| 415 | -0.250952 | 1.630752 | 0.770014 | 0.781814 | 3.422802 | 0.470543 | | -0.645932 | -0.956462 |
| 283 | 0.936914 | 1.255187 | 0.873409 | -1.288212 | -0.692891 | -0.202129 | | -0.926803 | 1.170732 |

116 rows × 8 columns

Next steps: [Generate code with X_test](#) [New interactive sheet](#)

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5, weights='distance', metric='euclidean')
knn.fit(X_train, y_train)
```

KNeighborsClassifier
 KNeighborsClassifier(metric='euclidean', weights='distance')

```
y_pred = knn.predict(X_test)
y_pred
```

```
array([1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1,
       1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1,
       1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
       0, 0, 1, 1, 1])
```

```
from sklearn.metrics import accuracy_score
print('Model accuracy score: {:.4f}'.format(accuracy_score(y_test, y_pred)))
```

```
Model accuracy score: 0.8190
```

```
y_pred_train = knn.predict(X_train)
print('Training-set accuracy score: {:.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

```
Training-set accuracy score: 1.0000
```

```
print('Training set score: {:.4f}'.format(knn.score(X_train, y_train)))
print('Test set score: {:.4f}'.format(knn.score(X_test, y_test)))
```

```
Training set score: 1.0000
Test set score: 0.8190
```

```
y_test.value_counts()
```

| Outcome | count |
|---------|-------|
| 0 | 78 |
| 1 | 38 |

```
dtype: int64
```

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print('Confusion matrix\n', cm)
print('\nTrue Negatives(TP) = ', cm[0,0])
print('\nTrue Positives(TN) = ', cm[1,1])
print('\nFalse Negatives(FP) = ', cm[0,1])
print('\nFalse Positives(FN) = ', cm[1,0])
```

```
Confusion matrix
```

```
[[69  9]
 [12 26]]
```

```
True Negatives(TP) = 69
```


```
True Positives(TN) = 26
```

```
False Negatives(FP) = 9
```

```
False Positives(FN) = 12
```

```
plt.figure(figsize=(6,4))
cm_matrix = pd.DataFrame(cm, columns=['Actual Benign:2', 'Actual Malignant:4'], index=['Predict Benign:2', 'Predict Malignant:4'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

<Axes: >



```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.88 | 0.87 | 78 |
| 1 | 0.74 | 0.68 | 0.71 | 38 |
| accuracy | | | 0.82 | 116 |
| macro avg | 0.80 | 0.78 | 0.79 | 116 |
| weighted avg | 0.82 | 0.82 | 0.82 | 116 |

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')
```